

Bayesian Personalized Ranking

Weike Pan

College of Computer Science and Software Engineering
Shenzhen University

Outline

- 1 Introduction
- 2 Method
- 3 Experiments
- 4 Conclusion
- 5 References

Recommendation with Implicit Feedback

- We may represent users' implicit feedback in a **matrix** form:

		...				
	?	1	?	?	?	?
	?	?	1	1	1	?
	?					1
	?					1
	?	1	?	1	?	1
	1	?	?	1	?	?

- If we can **estimate the missing values** (denoted as “?”) in the matrix or **rank the items directly**, we can make recommendations for each user.

Typical Steps in Recommendation with Implicit Feedback

For each user u :

- Step 1: Predict the preference of user u on item j , i.e., \hat{r}_{uj} , where $j \in \mathcal{I} \setminus \mathcal{I}_u$. We can use different methods, e.g.,
 - PopRank
 - User-based OCCF, item-based OCCF, hybrid OCCF
 - BPR
 - FISM
 - ...
- Step 2: Rank the items in $\mathcal{I} \setminus \mathcal{I}_u$ and use the top- k items with highest preference values to construct the recommendation list

Notations (1/2)

Table: Some notations.

n	user number
m	item number
$u \in \{1, 2, \dots, n\}$	user ID
$i, j \in \{1, 2, \dots, m\}$	item ID
$\mathcal{R} = \{(u, i)\}$	(user, item) pairs in training data
$y_{ui} \in \{1, 0\}$	indicator variable, $y_{ui} = 1$ if $(u, i) \in \mathcal{R}$
\mathcal{I}_u	preferred items by user u in training data
\mathcal{I}	the whole item set
\mathcal{U}	the whole user set

Notations (2/2)

Table: Some notations.

$b_i \in \mathbb{R}$	item bias
$d \in \mathbb{R}$	number of latent dimensions
$U_{u.} \in \mathbb{R}^{1 \times d}$	user-specific latent feature vector
$V_{i.} \in \mathbb{R}^{1 \times d}$	item-specific latent feature vector
\hat{r}_{ui}	predicted rating of user u on item i
T	iteration number in the algorithm

Pointwise Preference Assumption

The assumption of pointwise preference on an item [Hu et al., 2008, Pan et al., 2008] can be represented as follows,

$$\hat{r}_{ui} = 1, \hat{r}_{uj} = 0, i \in \mathcal{I}_u, j \in \mathcal{I} \setminus \mathcal{I}_u, \quad (1)$$

where 1 and 0 are used to denote “like” and “dislike” for an observed (user, item) pair and an unobserved (user, item) pair, respectively.

Notes:

- Treating all observed feedback as “likes” and unobserved feedback as “dislikes” may mislead the learning process.

Pairwise Preference Assumption

The assumption of pairwise preferences over two items [Rendle et al., 2009] relaxes the assumption of pointwise preferences, which can be represented as follows,

$$\hat{r}_{ui} > \hat{r}_{uj}, i \in \mathcal{I}_u, j \in \mathcal{I} \setminus \mathcal{I}_u \quad (2)$$

where the relationship $\hat{r}_{ui} > \hat{r}_{uj}$ means that **a user u is likely to prefer an item $i \in \mathcal{I}_u$ to an item $j \in \mathcal{I} \setminus \mathcal{I}_u$.**

Notes:

- Empirically, this assumption generates better recommendation results than the pointwise assumption.

Prediction Rule

The predicted rating of user u on item i ,

$$\hat{r}_{ui} = U_u \cdot V_i^T + b_i \quad (3)$$

Question:

- why not include b_u and μ

Likelihood of Pairwise Preferences (1/2)

The Bernoulli distribution of binary random variable $\delta((u, i) \succ (u, j))$ is defined as follows [Rendle et al., 2009],

$$\begin{aligned} \text{LPP}_u &= \prod_{i, j \in \mathcal{I}} Pr(\hat{r}_{ui} > \hat{r}_{uj})^{\delta((u, i) \succ (u, j))} [1 - Pr(\hat{r}_{ui} > \hat{r}_{uj})]^{[1 - \delta((u, i) \succ (u, j))]} \\ &= \prod_{(u, i) \succ (u, j)} Pr(\hat{r}_{ui} > \hat{r}_{uj}) \times \prod_{(u, i) \preceq (u, j)} [1 - Pr(\hat{r}_{ui} > \hat{r}_{uj})] \end{aligned}$$

where $(u, i) \succ (u, j)$ means that user u prefers item i to item j .

Likelihood of Pairwise Preferences (2/2)

We use $\sigma(\hat{r}_{uij})$ to approximate the probability $Pr(\hat{r}_{ui} > \hat{r}_{uj})$ [Rendle et al., 2009], where $\hat{r}_{uij} = \hat{r}_{ui} - \hat{r}_{uj}$, and have

$$\begin{aligned}
 \ln \text{LPP}_u &= \ln \prod_{(u,i) \succ (u,j)} \sigma(\hat{r}_{uij}) + \ln \prod_{(u,i) \preceq (u,j)} [1 - \sigma(\hat{r}_{uij})] \\
 &\approx \ln \prod_{(u,i) \succ (u,j)} \sigma(\hat{r}_{uij}) + \ln \prod_{(u,i) \succ (u,j)} [1 - \sigma(-\hat{r}_{uij})] \\
 &= \ln \prod_{(u,i) \succ (u,j)} \sigma(\hat{r}_{uij}) + \ln \prod_{(u,i) \succ (u,j)} \sigma(\hat{r}_{uij}) \\
 &= 2 \sum_{(u,i) \succ (u,j)} \ln \sigma(\hat{r}_{uij}) \\
 &= 2 \sum_{i \in \mathcal{I}_u} \sum_{j \in \mathcal{I} \setminus \mathcal{I}_u} \ln \sigma(\hat{r}_{uij})
 \end{aligned} \tag{4}$$

where $\sigma(z) = 1/(1 + e^{-z})$ is the sigmoid function.

Objective Function

Objective function,

$$\min_{\Theta} \sum_{u \in \mathcal{U}} \sum_{i \in \mathcal{I}_u} \sum_{j \in \mathcal{I} \setminus \mathcal{I}_u} f_{uij} \quad (5)$$

where

$f_{uij} = -\ln \sigma(\hat{r}_{uij}) + \frac{\alpha_u}{2} \|U_{u\cdot}\|^2 + \frac{\alpha_v}{2} \|V_{i\cdot}\|^2 + \frac{\alpha_v}{2} \|V_{j\cdot}\|^2 + \frac{\beta_v}{2} \|b_i\|^2 + \frac{\beta_v}{2} \|b_j\|^2$
 and $\Theta = \{U_{u\cdot}, u = 1, 2, \dots, n; V_{i\cdot}, b_i, i = 1, 2, \dots, m\}$ denotes the set of parameters to be learned.

Gradients

For a randomly sampled triple (u, i, j) , we have the gradients,

$$\nabla U_u = \frac{\partial f_{uij}}{\partial U_u} = -\sigma(-\hat{r}_{uij})(V_i - V_j) + \alpha_u U_u, \quad (6)$$

$$\nabla V_i = \frac{\partial f_{uij}}{\partial V_i} = -\sigma(-\hat{r}_{uij})U_u + \alpha_v V_i, \quad (7)$$

$$\nabla V_j = \frac{\partial f_{uij}}{\partial V_j} = -\sigma(-\hat{r}_{uij})(-U_u) + \alpha_v V_j, \quad (8)$$

$$\nabla b_i = \frac{\partial f_{uij}}{\partial b_i} = -\sigma(-\hat{r}_{uij}) + \beta_v b_i, \quad (9)$$

$$\nabla b_j = \frac{\partial f_{uij}}{\partial b_j} = -\sigma(-\hat{r}_{uij})(-1) + \beta_v b_j, \quad (10)$$

where $\hat{r}_{uij} = \hat{r}_{ui} - \hat{r}_{uj}$.

Update Rules

For a randomly sampled triple (u, i, j) , we have the update rules,

$$U_{u.} = U_{u.} - \gamma \nabla U_{u.}, \quad (11)$$

$$V_{i.} = V_{i.} - \gamma \nabla V_{i.}, \quad (12)$$

$$V_{j.} = V_{j.} - \gamma \nabla V_{j.}, \quad (13)$$

$$b_i = b_i - \gamma \nabla b_i, \quad (14)$$

$$b_j = b_j - \gamma \nabla b_j, \quad (15)$$

where γ is the learning rate.

Algorithm

```
1: Initialize the model parameters  $\Theta$ 
2: for  $t = 1, \dots, T$  do
3:   for  $t_2 = 1, \dots, |\mathcal{R}|$  do
4:     Randomly pick up a pair  $(u, i) \in \mathcal{R}$ 
5:     Randomly pick up an item  $j$  from  $\mathcal{I} \setminus \mathcal{I}_u$ 
6:     Calculate the gradients via Eq.(6-10)
7:     Update the model parameters via Eq.(11-15)
8:   end for
9: end for
```

Figure: The SGD algorithm for BPR.

Data Set

- We use the files `u1.base` and `u1.test` of MovieLens100K¹ as our training data and test data, respectively.
- user number: $n = 943$; item number: $m = 1682$.
- `u1.base` (training data): 80000 rating records, and the density (or sparsity) is $80000/943/1682 = 5.04\%$.
- `u1.test` (test data): 20000 rating records.
- **Pre-processing (for simulation)**: we only keep the (user, item) pairs with ratings 4 or 5 in `u1.base` and `u1.test` as preferred (user, item) pairs, and remove all other records. Finally, we obtain **`u1.base.OCCF`** and **`u1.test.OCCF`**.

¹<http://grouplens.org/datasets/>

Evaluation Metrics

- *Pre@5*: The precision of user u is defined as,

$$Pre_u@k = \frac{1}{k} \sum_{\ell=1}^k \delta(i(\ell) \in \mathcal{I}_u^{te}),$$

where $\delta(x) = 1$ if x is true and $\delta(x) = 0$ otherwise. Then, we have $Pre@k = \sum_{u \in \mathcal{U}^{te}} Pre_u@k / |\mathcal{U}^{te}|$.

- *Rec@5*: The recall of user u is defined as,

$$Rec_u@k = \frac{1}{|\mathcal{I}_u^{te}|} \sum_{\ell=1}^k \delta(i(\ell) \in \mathcal{I}_u^{te}),$$

which means how many preferred items are recommended in the top- k list. Then, we have $Rec@k = \sum_{u \in \mathcal{U}^{te}} Rec_u@k / |\mathcal{U}^{te}|$.

Initialization of Model Parameters

We use the statistics of training data to initialize the model parameters,

$$b_i = \left(\frac{1}{n} \sum_{u=1}^n y_{ui} \right) - \mu$$

$$V_{ik} = (r - 0.5) \times 0.01, k = 1, \dots, d$$

$$U_{uk} = (r - 0.5) \times 0.01, k = 1, \dots, d$$

where r ($0 \leq r < 1$) is a random variable, and $\mu = \sum_{u=1}^n \sum_{i=1}^m y_{ui} / n / m$.

Parameter Configurations

We fix $\gamma = 0.01$, and search the best values of the following parameters,

- $\alpha_u = \alpha_v = \beta_v \in \{0.001, 0.01, 0.1\}$
- $T \in \{100, 500, 1000\}$
- $d = 20$

Finally, we use $\gamma = 0.01$, $\alpha_u = \alpha_v = \beta_v = 0.01$, $T = 500$ and $d = 20$.

Results

Table: Prediction performance of PopRank and BPR on MovieLens100K (u1.base.OCCF, u1.test.OCCF). Note that the time cost using Java in my PC is less than 15 seconds.

	PopRank	BPR
<i>Pre@5</i>	0.2338	0.3864
<i>Rec@5</i>	0.0571	0.1184

Conclusion

- The pairwise preference assumption is useful.

Homework

- Read the code of BPR implementation in MyMediaLite²
 - Pay attention to different sampling strategies
- Implement BPR and conduct empirical studies on u2.base.OCCF, u2.test.OCCF of MovieLens100K with similar pre-processing
- Read the UAI 2009 paper [Rendle et al., 2009]

²<http://www.mymedialite.net/>



Hu, Y., Koren, Y., and Volinsky, C. (2008).

Collaborative filtering for implicit feedback datasets.

In *Proceedings of the 2008 Eighth IEEE International Conference on Data Mining, ICDM '08*, pages 263–272, Washington, DC, USA. IEEE Computer Society.



Pan, R., Zhou, Y., Cao, B., Liu, N. N., Lukose, R., Scholz, M., and Yang, Q. (2008).

One-class collaborative filtering.

In *Proceedings of the 2008 Eighth IEEE International Conference on Data Mining, ICDM '08*, pages 502–511, Washington, DC, USA. IEEE Computer Society.



Rendle, S., Freudenthaler, C., Gantner, Z., and Schmidt-Thieme, L. (2009).

Bpr: Bayesian personalized ranking from implicit feedback.

In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence, UAI '09*, pages 452–461, Arlington, Virginia, United States. AUAI Press.